

Programming Assignment 1: A Simple Multi-threaded HTTP Server

Due: September 19th, 2017 23:59:59pm

In this project, you will implement a simple multi-threaded HTTP server that only accepts HTTP GET requests and returns the desired content to the client. This project is worth **7%** of your total score in this course. You **must** work **by yourself** on this project and use **C/C++, Java, or Python** to implement this project.

1 Background

HTTP GET requests are used to request data from the specified source. For example, if we want to access a webpage at `http://www.foo.com/bar.html`, our browser will send something similar to the following HTTP GET request.

```
GET /bar.html HTTP/1.1
Host: www.foo.com
User-Agent: Mozilla/5.0 (X11; Ubuntu; Linux x86_64; rv:28.0) Gecko/20100101 Firefox/28.0
Accept: text/html,application/xhtml+xml,application/xml;q=0.9,*/*;q=0.8
Accept-Language: en-US,en;q=0.5
Accept-Encoding: gzip, deflate
Connection: keep-alive
```

In the above example, the Firefox browser is the client, and `Host: www.foo.com` is the server host name. The client is requesting server resource located at `/bar.html`. Usually, HTTP servers run on port 80. So, by default, the server port number is 80. If, for example, the HTTP server runs on port 8080, then the URL would be `http://www.foo.com:8080/bar.html`, and the `Host` field in the HTTP request will also contain the port number: `www.foo.com:8080`.

The HTTP server replies with HTTP response. A valid HTTP response includes: (1) the status line, (2) the response header, and (3) the message body. For example, the text below shows an HTTP response. Note that there is an empty line between the response header and the message body. This empty line indicates the end of header fields¹.

```
HTTP/1.1 200 OK
Date: Thu, 02 Apr 2015 01:51:49 GMT
Server: Apache/2.2.16 (Debian)
Last-Modified: Tue, 10 Feb 2015 17:56:15 GMT
Accept-Ranges: bytes
Content-Length: 2693
Content-Type: text/html
```

```
... <content of the bar.html>
```

¹<https://www.w3.org/Protocols/rfc2616/rfc2616-sec4.html>

2 Multi-threaded HTTP Server

Your HTTP server should support a subset of the HTTP standard. It only needs to serve HTTP GET requests and should only support basic portions of the standard. That is, you do not need to support persistent connections, request pipelining or other, advanced portions of the standard. Your server must only correctly handle HTTP requests where a single request/response pair is sent over a single TCP connection. Specifically, your server should operate as follows:

- Your HTTP server should not take any command line arguments. It should create a TCP server socket to listen for incoming TCP connections on an unused port, and output the host name and port number the HTTP server is running on. **Note: You must use sockets in this project. Any implementation without the use of sockets will receive zero points.**
- It should look for a directory called `www`, **located in the same parent directory as your HTTP server executable**. The `www` directory should contain resources you want your HTTP server to serve. If this directory does not exist, the HTTP server should output an error message and quit.
- When a new request comes in, the HTTP server accepts the connection, establishing a TCP connection with the client.
- The HTTP server parses the HTTP request from the client and prepares an HTTP response.
- The HTTP server looks up the requested resource from its local `www` directory.
 - If the requested resource is located by the HTTP server, prepare a response with status code “200 OK”, followed by response header and the content of the requested resource. The response header should include the following fields:
 - * `Date` The date and time the response is originated in the format defined by RFC 7231 Date/Time Formats ².
 - * `Server` A name for your HTTP server. You can be very creative here.
 - * `Last-Modified` Last modified time and date of the requested resource, also in RFC 7231 Date/Time Formats.
 - * `Content-Type` The MIME type of this content. You can use the `mime.types` file located at `/etc/mime.types` to determine the correct MIME type given a filename extension. If a filename extension is not found in the `mime.types` file, you can use `application/octet-stream`.
 - * `Content-Length` The length of the requested resource in bytes.
 - If the requested resource does not exist, reply with status code “404 Not Found”. You can be creative and write a fancy custom error page to return to the client.
- After finishing sending the response, the HTTP server closes the connection socket to the client.

To maintain a high level of throughput even under multiple simultaneous client requests, different requests should be handled in different threads, with a new thread created for each client request.

2.1 How to test your implementation

You **must** test your HTTP server on the remote nodes of CS department computers: `remote.cs.binghamton.edu`. When accessing `remote.cs.binghamton.edu`, you are actually redirected to one of the 8 REMOTE machines

²<https://tools.ietf.org/html/rfc7231#section-7.1.1.1>

({remote00, remote01, ..., remote06, remote07}.cs.binghamton.edu) using DNS redirection. So to test your implementation, you need to run your HTTP server on one of the REMOTE nodes and set your web client, e.g., `wget`³, to use the correct server host name: “remoteXX.cs.binghamton.edu”, instead of “remote.cs.binghamton.edu”.

For example, suppose your HTTP server is started on `remote02.cs.binghamton.edu` port 47590. You can run the following command in an **EMPTY** directory on a different REMOTE computer to download the resource:

```
wget http://remote02.cs.binghamton.edu:47590/bar.html
```

Note:

- You **MUST** replace `bar.html` with a valid resource in your HTTP server directory!
- Since this is only an HTTP server, not an HTTPS server, your URL **MUST** use HTTP, not HTTPS.
- You may also run your HTTP server on your own machine and use Wireshark⁴ for debugging. However, you must test your implementation on REMOTE computers before submission.
- Your web browser has a cache. A second request for the same URL may be directly served by your browser cache without going through your server. Therefore, you are recommended to use `wget` for debugging and testing.

After `wget` successfully downloads the requested resource, use the `diff`⁵ command to check if the downloaded resource matches the resource located in the `www` directory.

`wget` has a few options that can help you debug your implementation. For example, the `--limit-rate` option allows you limit the download speed to a pre-specified amount. This can be helpful if you want to debug your multi-threaded implementation. If you need to inspect if your HTTP server constructs the HTTP response header correctly, you can take advantage of the `--save-headers` option. More information can be found at the manual page of `wget`.

Upon successfully serving a request, your HTTP server must write to standard output (stdout) the following items:

requested resource

client IP client’s IP address in dotted decimal representation

client port client’s port number

access times the number of times this resource has been requested since the start of the HTTP server

These four items of a same request should be printed to a same line and separated by the “|” character. Below is an example:

```
/bar.html|128.226.118.20|4759|1
/pics/foo.jpg|128.226.118.20|6001|1
/bar.html|128.226.118.26|6550|2
/bar.html|128.226.118.21|10039|3
/pics/foo.jpg|128.226.118.21|5300|2
/pics/foo.jpg|128.226.118.26|36500|3
```

³<http://linux.die.net/man/1/wget>

⁴<https://www.wireshark.org/>

⁵<http://linux.die.net/man/1/diff>

3 How to submit

To submit the project, you should first create a directory whose name is identical to your BU email ID. For example, if your email ID is `jdoue@binghamton.edu`, you should create a directory called `jdoue-pl`.

You should put the following files into this directory:

1. All of your source code that implements the HTTP server. Note that your server executable should not take any command line arguments and must output the host name and port number it is running on.
2. A `Makefile` to compile your source code into one executable, which should be named `server`.
3. A `Readme` file describing how to compile and run your code on `remote.cs.binghamton.edu` computers, a brief description of your implementation, and sample input/output.
4. A `STATEMENT` file, containing the following statement followed by the student's full name:

"I have done this assignment completely on my own. I have not copied it, nor have I given my solution to anyone else. I understand that if I am involved in plagiarism or cheating I will have to sign an official form that I have cheated and that this form will be stored in my official university record. I also understand that I will receive a grade of **0** for the involved assignment and my grade will be reduced by one level (e.g., from A to A- or from B+ to B) for my first offense, and that I will receive a grade of **"F" for the course** for any additional offense of any kind."

Compress the directory (e.g., `tar czvf jdoue-pl.tgz jdoue-pl`) and submit the tarball (in `tar.gz` or `tgz` format) to myCourses. Again, if your email ID is `jdoue@binghamton.edu`, you should name your submission: `jdoue-pl.tar.gz` or `jdoue-pl.tgz`. Failure to use the right file name will result in a 5% point deduction.

Your project will be graded on the CS Department computers `remote.cs.binghamton.edu`. If your code does not compile on or cannot correctly run on the CS computers, you will receive zero points.

Your project must be your original work. We will use MOSS⁶ to detect plagiarism in the projects.

Appendix

You can download the following files into your local `www` directory and have your HTTP server serve these resources.

```
http://www.cs.binghamton.edu/~yaoliu/courses/cs557/test.html
http://www.cs.binghamton.edu/~yaoliu/courses/cs557/pdf-sample.pdf
http://www.cs.binghamton.edu/~yaoliu/courses/cs557/lena_std.tif
http://www.cs.binghamton.edu/~yaoliu/courses/cs557/skype-ubuntu-precise_4.3.0.37-1_i386.deb
```

⁶<https://theory.stanford.edu/~aiken/moss/>